# Matrix Inversion Optimization

# Angus Morrison, V00937482, Software Engineering Daisy Kwan, V00919518, Computer Science

Department of CS/ECE, University of Victoria SENG440 Embedded Systems angusmorrison@uvic.ca daisykwan@uvic.ca

# **Constraints and Specifications**

\_\_\_\_

#### Specs:

- 10x10 matrix
- 12 bit width elements
- Use fix point arithmetic
- Use real value matrices
- Condition Number threshold < 9

#### We want to:

- Vectorize the code

#### Processor:

- Arm
- Qemu emulator on SENG440 sever
- Supports SIMD (NEON) intrinsics

# History

Carl Friedrich Gauss developed Gaussian elimination

- Gauss-Jordan elimination initially described by B.-I. Clasen 1888

Wilhelm Jordan expanded on Gaussian elimination in the same

year



# Background

- Gauss Jordan elimination
  - The use of three elementary row operations to transform a matrix into reduced row-echelon form
  - A matrix is in reduced row-echelon form when all four conditions are met

$$\begin{bmatrix} 1 & 2 & 3 \\ 3 & 2 & 1 \\ 2 & 1 & 2 \end{bmatrix} \sim \dots \sim \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

# Possible Application of Interest

- Solving systems of equations

- Communications
  - 12 bits is commonly used in communication

- Computer graphics
  - Used in 3D renderings, ray casting, transformations, and physical simulations

# **Converting From Float to Integer**

Bit Width: 12

- 1 bit sign
- 11 bits magnitude

Scale Factor: 2<sup>11</sup>

```
\begin{bmatrix} 0.99 & 0.92 & -0.86 \\ -0.56 & 0.32 & 0.53 \\ -0.63 & -0.88 & 0.29 \end{bmatrix} - -0.99 \times 2^{11} = 0.92 \times 2^{11}
                                                                                                          -0.86 \times 2^{11}
    \begin{vmatrix} -0.56 \times 2^{11} & 0.32 \times 2^{11} & 0.53 \times 2^{11} \\ -0.63 \times 2^{11} & -0.88 \times 2^{11} & 0.29 \times 2^{11} \end{vmatrix} - 
   egin{bmatrix} 2028 & 1884 & -1761 \\ -1147 & 655 & 1085 \\ -1290 & -1802 & 593 \end{bmatrix}
```

# Steps taken 1 - Hard coding the identity matrix

#### Original Code

```
int main(int argc, char *argv[]) {
    float matrix[10][20] = {
        {10, 1, 1, 1, 1, 1, 1, 1, 1, 1},
        {1, 9, 1, 1, 1, 1, 1, 1, 1, 1},
        {1, 1, 8, 1, 1, 1, 1, 1, 1, 1},
        {1, 1, 1, 7, 1, 1, 1, 1, 1, 1},
        {1, 1, 1, 1, 6, 1, 1, 1, 1, 1},
        {1, 1, 1, 1, 1, 5, 1, 1, 1, 1},
        {1, 1, 1, 1, 1, 1, 4, 1, 1, 1},
        {1, 1, 1, 1, 1, 1, 1, 3, 1, 1},
        {1, 1, 1, 1, 1, 1, 1, 1, 2, 1},
        {1, 1, 1, 1, 1, 1, 1, 1, 1, 1},
    printf("Starting the inversion\n"):
    // Do the inversion
    /* Augmenting Identity Matrix of Order 10 */
    for (int i = 0; i < 10; i++) {
        for (int j = 0; j < 10; j++) {
            // if i == j, we are on the diagonal, so make the identity matrix
            if (i == j) {
                matrix[i][j + 10] = 1;
            } else {
                matrix[i][j + 10] = 0;
```

#### Gauss-Jordan Code

# Steps taken 2 - Register keyword

```
Original
for (int i = 0; i < 10; i++) {
   for (int j = 0; j < 10; j++) {</pre>
```

```
register int i = 0;
register int j = 0;
register int k = 0;
```

# Steps taken 2 - Register keyword (Assembly)

#### **Original**

```
.L13:
     1dr
          r3, [fp, #-60]
     add
         r3, r3, #1
          r3, [fp, #-60]
     str
.L11:
         r3, [fp, #-60]
     ldr
           r3, #9
     CMP
     ble
           .L14
     ldr r3, [fp, #-64]
     add r3, r3, #1
           r3, [fp, #-64]
```

#### Optimized

```
ldr r1, [fp, #-1896]
sub r1, r1, #1
str r1, [fp, #-1896]
.L23:
ldr r2, [fp, #-1896]
cmp r2, #0
bge .L24
ldr r3, [fp, #-1900]
sub r3, r3, #1
str r3, [fp, #-1900]
```

This shows in two places that using the register keyword didn't do anything as the program still loads the value for the counters decrements and then stores the value back in memory

# Steps taken 3 - Loop decrement

#### Original

# // Devide each row by the diagnal to get the inverse for (int i = 0; i < 10; i++) { float diagonal = matrix[i][i]; for (int j = 0; j < 20; j++) { matrix[i][j] /= diagonal; }</pre>

```
// Divide each row by the diagonal to get the inverse
// move diagnal to a register and so that it is not reinitialized each loop
register float diagonal;
for (i = 9; i >= 0; i--){
    diagonal = matrix[i][i];
    for (j = 19; j >= 0; j--) {
        matrix[i][j] /= diagonal;
    }
}
```

# Steps taken 3 - Loop decrement (Assembly )

```
.L23:
     1dr
           r2, [fp, #-1896]
           r2, #0
     cmp
           .L24
     bge
           r3, [fp, #-1900]
     1dr
           r3, r3, #1
     sub
           r3, [fp, #-1900]
     str
.L22:
           r4, [fp, #-1900]
     ldr
           r4, #0
      cmp
            .L25
     bge
```

```
.L25:

mov r1, #19

str r1, [fp, #-1896]

b .L23
```

This shows that the loop decrement didn't change much in the code because the register keyword didn't work. Due to the load and store we weren't able to check the zero flag to speed up the code and instead we still have to compare to zero

# Steps taken 4 - Loop removal

#### **Original**

```
print+("Starting the inversion\n");
// Do the inversion
/* Augmenting Identity Matrix of Order 10 */
for (int i = 0; i < 10; i++) {
    for (int i = 0; i < 10; i++) {
       // if i == j, we are on the diagonal, so make the identity matrix
       if (i == j) {
            matrix[i][j + 10] = 1;
        } else {
            matrix[i][j + 10] = 0;
// Interchange the rows of the matrix, starting from the last row
for (int i = 9; i > 0; i--) {
    // Swap rows
   if (matrix[i - 1][0] < matrix[i][0]) {
        float temp[20]:
        for (int j = 0; j < 20; j++) {
            temp[j] = matrix[i][j];
            matrix[i][j] = matrix[i - 1][j];
            matrix[i - 1][j] = temp[j];
```

```
float matrix[10][20] = {
    {10, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0},
    {1, 9, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0},
    {1, 1, 8, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0},
    \{1, 1, 1, 7, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0\}
    \{1, 1, 1, 1, 6, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0\}
    {1, 1, 1, 1, 1, 5, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0},
    {1, 1, 1, 1, 1, 1, 4, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0},
    \{1, 1, 1, 1, 1, 1, 1, 1, 3, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0\},\
    {1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0},
    {1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1},
register int i = 0;
register int i = 0:
register int k = 0:
// printf("Apply Gauss-Jordan elimination\n");
// Interchange the rows of the matrix, starting from the last row
// create local variable so that memory only has to be accessed once per loop
// I don't think we can use the register keyword as we would run out of registers to use
// I removed the nested loops and replaced them with memcpy
float array_temp[20];
for(i = 8: i>0: i--){
    if(matrix[i][0] < matrix[i+1][0]){
        memcpy(array_temp, matrix[i+1], 20*sizeof(int));
        memcpy(matrix[i+1], matrix[i], 20*sizeof(int));
        memcpy(matrix[i], array temp, 20*sizeof(int));
```

# **Steps taken 5 - Vectorization**

#### Original

```
for (i = 9; i >= 0; i--) {
if(int matrix[i][i] == 0){
   printf("dividing by 0 at i=%d\n", i);
   return 0;
   for (j = 9; j >=0; j--) {
        if (i != i) {
   temp = (int matrix[j][i] * 2048)/ int matrix[i][i]; // Have to ensure the scale factor remains
   for(k=0; k<20; k = k+4){ // runs 5 times
        // load the vector values
       i_vector_row = vld1q_u32(int_matrix[i] + k);
        j vector row = vld1q u32(int matrix[j] + k);
       // Do the multiplication
        i vector row = vmulq n u32(i vector row, temp);
       i_vector_row = vmulq_n_u32(i_vector_row, 1/2048); // have to ensure the scale factor remains
        j vector row = vsubq u32(j vector row, i vector row);
       vst1q u32(int matrix[j] + k, j vector row);
```

# Using memcpy() - Memory accesses

```
for (int i = 9; i > 0; i--) {
    // Swap rows
    if (matrix[i - 1][0] < matrix[i][0]) {
        float temp[20];
        for (int j = 0; j < 20; j++) {
            temp[j] = matrix[i][j];
            matrix[i][j] = matrix[i - 1][j];
            matrix[i - 1][j] = temp[j];
        }
    }
}</pre>
```

```
int array_temp[20];
for(i = 8; i>0; i--){
    if(int_matrix[i][0] < int_matrix[i+1][0]){
        memcpy(array_temp, int_matrix[i+1], 20*sizeof(int));
        memcpy(int_matrix[i+1], int_matrix[i], 20*sizeof(int));
        memcpy(int_matrix[i], array_temp, 20*sizeof(int));
    }
}</pre>
```

Using memcpy didn't reduce the number of memory access operations from the original to the optimized code however it did allow us to completely remove a loop freeing a register and reducing the number of operations

### Results

- The optimization increased the number lines in the assembly file from 854 to 871
- This was completely different than what we were expecting the result to be
- There was no change in program runtime to the 1000th of a second
- There was no change in runtime with ill and well conditioned matrices

# **Results Unoptimized**

```
Performance counter stats for './angus_unoptimized.exe':
            1.06 msec task-clock:u
                                                   0.133 CPUs utilized
                      context-switches:u
                                                   0.000 K/sec
                      cpu-migrations:u
                                                   0.000 K/sec
             342
                      page-faults:u
                                                   0.322 M/sec
       1,422,046
                      cycles:u
                                                 1.338 GHz
       1,891,280
                                               # 1.33 insn per cycle
                      instructions:u
         368,907
                      branches:u
                                               # 347.009 M/sec
                      branch-misses:u
                                                    3.09% of all branches
          11,396
     0.007993440 seconds time elapsed
     0.001003000 seconds user
     0.001003000 seconds sys
```

# **Results Optimized**

```
Performance counter stats for './angusmorrison.exe':
                                             # 0.146 CPUs utilized
           1.11 msec task-clock:u
                     context-switches:u
                                             # 0.000 K/sec
                     cpu-migrations:u
                                                  0.000 K/sec
             341
                     page-faults:u
                                                  0.308 M/sec
       1,407,980
                     cycles:u
                                             # 1.270 GHz
       1,890,799
                    instructions:u
                                             # 1.34 insh per cycle
         368,818
                    branches:u
                                             # 332.729 M/sec
                     branch-misses:u
                                                  3.05% of all branches
          11,248
     0.007566833 seconds time elapsed
     0.000571000 seconds user
     0.001714000 seconds sys
```

## **Conclusion**

We do not recommend trying to optimize matrix inversion code manually. The amount of coding and debugging time that went in for an unnoticeable change in runtime does not justify the costs.

Additionally there was no noticeable change in runtime between a ill and well conditioned matrix. This leads us to believe that the system is more robust albeit slow.

### References

Brilliant. (n.d.). Gauss-Jordan Elimination. Brilliant Math & Science Wiki.

https://brilliant.org/wiki/gaussian-elimination/#:~:text=Carl%20Friedrich%20Gauss%20championed%20the,inverses%2C%20Gauss%2DJordan%20elimination.

Knill, O. (n.d.). *About Gauss-Jordan elimination*. Math 22a Harvard College Fall 2018. https://people.math.harvard.edu/~knill/teaching/math22a2018/exhibits/gaussjordan/index.html

M.7 Gauss-Jordan Elimination: Stat Online. PennState: Statistics Online Courses. (n.d.).

https://online.stat.psu.edu/statprogram/reviews/matrix-algebra/gauss-jordan-elimination#:~:text=The%20purpose%20of%20Gauss%2DJordan,the%20bottom%20of%20the%20matrix